

VectorForge

GPU-native vector database for real-time RAG at inference scale

Team

- CEO - ex-Staff Engineer, Pinecone query engine (3 yrs)
- CTO - PhD, approximate nearest-neighbour search, 400+ citations
- CTO maintains annlib, 6,100 GitHub stars

Traction & Revenue

- 3 enterprise design partners in production
- 220M vectors served per day
- 3x QoQ query growth
- \$18K MRR from two paid design partners
- Third pilot converting next quarter